

Verifying an agentic GIS routing workflow


Fiber route comparison on the shortened Mombasa corridor: baseline, roads, power and rail

Justin Rogers

GG5 662: Agentic GeoAI · Assignment 2

September 28, 2026

Week 4 audit: shortened corridor routes, connectors and flagged crossings



Route and infrastructure data © OpenStreetMap contributors (ODbL). Basemap: Esri World Light Gray Canvas (Esri, HERE, Garmin, © OpenStreetMap contributors).

Map data © OpenStreetMap contributors (ODbL); basemap Esri

Verifying an agentic GIS routing workflow


2026-09-27

About 15 s. Last week an agent built a fiber-routing workflow for the shortened Mombasa corridor and produced routes, a map and a cost table. This week the question is not whether it ran, but whether the results are correct and defensible. I will give the agent framework briefly, then spend most of the time on the checks and what they found.

Verifying an agentic GIS routing workflow

Fiber route comparison on the shortened Mombasa corridor: baseline, roads, power and rail

Justin Rogers
GG5 662: Agentic GeoAI · Assignment 2
September 28, 2026



Agent framework

Goal

Compare baseline, roads, power and rail routes between the exact route1_short endpoints (east = start); report length and USD 10/m cost.

Agency

Agent cleans, nodes, repairs gaps and routes. **Human** fixed endpoints, 50 m gap limit, 500 m endpoint limit and the rate.

Environment

Saved o3_01 OpenStreetMap (OSM) GeoPackages, AGENTS.md rules, 3,000 m study rectangle, measured in EPSG:32737. No new downloads.

Feedback

Stop with endpoint_too_far, no_path or empty_network. A failed check means revise, never relax a threshold.

Tools

geopandas, shapely, networkx (Dijkstra), pandas; matplotlib + contextily for maps.

Verification

Independent audit script, synthetic cases, controlled reruns. Real-world claims need imagery or authoritative data.

2026-09-27

Verifying an agentic GIS routing workflow

Agent framework

About 40 s. Goal: four candidate routes between the exact endpoints of the shortened baseline, easterly point as start, with length and a classroom cost of 10 dollars per metre. Environment: only the saved o3_01 files and the AGENTS.md rules. Tools: geopandas, shapely, networkx. Agency: the agent chooses how to clean, node, repair and route, but a human set the endpoints, the 50 m gap limit, the 500 m endpoint limit and the rate. Feedback: stop with a named status rather than invent a route; a failed check means revise, not raise a threshold. Verification: an audit that recomputes everything from saved files, plus outside evidence for real-world claims.

Agent framework		
Goal Compare baseline, roads, power and rail routes between the exact routes_short endpoints (east = start); report length and USD 10/m cost.	Environment Saved o3_01 OpenStreetMap (OSM) GeoPackages, AGENTS.md rules, 3,000 m study rectangle, measured in EPSG:32737. No new downloads.	Tools geopandas, shapely, networkx (Dijkstra), pandas; matplotlib + contextily for maps.
Agency Agent cleans, nodes, repairs gaps and routes. Human fixed endpoints, 50 m gap limit, 500m endpoint limit and the rate.	Feedback Stop with endpoint_too_far, no_path or empty_network. A failed check means revise, never relax a threshold.	Verification Independent audit script, synthetic cases, controlled reruns. Real-world claims need imagery or authoritative data.

Where the human decides and where the agent revises

Stays with a human

Whether an inferred connector or a 2D crossing is actually buildable, and which corridor to pursue.

Feedback that changes the next action

The audit flags the rail route switching lines at a bridge (−3.99374, 39.54626) with no shared OSM node.

Imagery showed an overpass, so the agent should **remove that transfer and reroute rail**, then report the new length or no_path. It should not raise a threshold to keep the old answer.

2026-09-27


Verifying an agentic GIS routing workflow

Where the human decides and where the agent revises

About 25 s. One decision that must stay human: whether anything the model inferred, a connector or a crossing, can actually be built, and which corridor to pursue; that depends on access rights and engineering that are not in the data. One piece of feedback that should change the agent’s next action: the audit found the rail route switching onto a bridge where the two lines share no OSM node. Imagery confirmed they are at different levels, an overpass, so the agent should delete that transfer and reroute, accepting a longer route or no_path, rather than loosening a threshold.

Stays with a human	Feedback that changes the next action
Whether an inferred connector or a 2D crossing is actually buildable, and which corridor to pursue.	The audit flags the rail route switching lines at a bridge (−3.99374, 39.54626) with no shared OSM node. Imagery showed an overpass, so the agent should remove that transfer and reroute rail , then report the new length or no_path. It should not raise a threshold to keep the old answer.

Approach: checks designed around the agent, not after it



Independent of the code. The audit never imports the routing script; it recomputes every value from the saved files and the 03_01 inputs.

Not independent of the author. The same AI tool wrote both scripts, so they can share a wrong assumption, such as treating every 2D crossing as a junction. The audit can only confirm the computations, so the flagged sites were also checked in satellite imagery, which is independent of both OSM and the code.

2026-09-27

Verifying an agentic GIS routing workflow


Approach: checks designed around the agent, not after it



About 25 s. I followed the 04_02 validation stack: data QA, the routing agent, then a separate audit script that imports no routing code and rebuilds every length, cost and connector from the saved files; controlled reruns that change only the gap tolerance; outside evidence at the sites the audit flags; and a human applies the labels. Caveat: the same AI tool wrote both scripts, so the audit is verification, not independent validation.

Route map: what the agent built and what it inferred

Week 4 audit: shortened corridor routes, connectors and flagged crossings



Route and infrastructure data © OpenStreetMap contributors (ODbL). Basemap: Esri World Light Gray Canvas (Esri, HERE, Garmin, © OpenStreetMap contributors).

Green rings: inferred connectors. Roads 0.59 m and 0.72 m; rail 194.8 m and 372.7 m. All four are endpoint connectors: hypothetical new construction.

Red cross: flagged crossing. The rail route switches onto a bridge-tagged line where the two lines share no OSM vertex.

Lengths. Baseline 22,896.8 m; roads 22,897.6 m; rail 25,902.4 m. Power: endpoint_too_far.

Data © OpenStreetMap contributors (ODbL). Basemap: Esri World Light Gray Canvas.

Verifying an agentic GIS routing workflow

2026-09-27

Route map: what the agent built and what it inferred

About 35 s. Grey baseline and black roads route overlap almost everywhere, under one metre apart in total length. Purple is rail, 3 km longer. Green rings mark every place the model invented geometry: two sub-metre snaps for roads, and 195 m and 373 m links for rail, all endpoint connectors, drawn separately from mapped infrastructure. The red cross is where the rail route jumps onto a bridge at a point with no shared node, the one failed check. Power has no route: its start is 1.8 km from any mapped power line.

Route map: what the agent built and what it inferred



Green rings: inferred connectors. Roads 0.59 m and 0.72 m; rail 194.8 m and 372.7 m. All four are endpoint connectors: hypothetical new construction.

Red cross: flagged crossing. The rail route switches onto a bridge-tagged line where the two lines share no OSM vertex.

Lengths. Baseline 22,896.8 m; roads 22,897.6 m; rail 25,902.4 m. Power: endpoint_too_far.

Data © OpenStreetMap contributors (ODbL). Basemap: Esri World Light Gray Canvas.

Verification evidence: 75 reproducible checks

73

PASS

2

FAIL: rail crossing (audit + imagery)

0

NOT_CHECKED

Category	PASS	FAIL	NOT_CHK	What it checks
Provenance and data QA	7	0	0	counts, CRS, invalid geometry, duplicates
Comparison table and cost	9	0	0	four rows, blank unavailable values, length × USD 10/m
Route geometry and length	23	0	0	endpoints, continuity, metres, edge sum within 0.01 m
Connectors	6	0	0	limits, containment, counted once
Unavailable route (power)	1	0	0	endpoint distance recomputed
Synthetic cases	20	0	0	6 cases × 3 runs, plus the 230 m oracle
Sensitivity	5	0	0	only the gap setting changed
Spatial assumption	0	1	0	route transfers only at OSM junctions
External evidence	2	1	0	satellite imagery at the 3 flagged sites

Verifying an agentic GIS routing workflow

2026-09-27

Verification evidence: 75 reproducible checks

About 35 s. Each check records method, expected, observed, evidence file and PASS, FAIL or NOT_CHECKED; missing evidence never counts as a pass. The computational checks all pass: exact endpoints, geometry length equal to the traversed-edge sum within 1 cm, metres not degrees (UTM within 0.036% of geodesic), cost equals length times 10, connectors counted once. Synthetic cases cover the five AGENTS.md situations plus connector-counted-once in all three runs, and a hand oracle: 80 + 120 + 30 m gap = 230 m and USD 2,300. The two fails are the interesting ones, and they are the same place: the audit flagged the rail crossing, and the imagery confirmed it is an overpass.

Verification evidence: 75 reproducible checks				
	73 PASS	2 FAIL: rail crossing (audit + imagery)	0 NOT_CHECKED	
Category	PASS	FAIL	NOT_CHK	What it checks
Provenance and data QA	7	0	0	counts, CRS, invalid geometry, duplicates
Comparison table and cost	9	0	0	four rows, blank unavailable values, length × USD 10/m
Route geometry and length	23	0	0	endpoints, continuity, metres, edge sum within 0.01 m
Connectors	6	0	0	limits, containment, counted once
Unavailable route (power)	1	0	0	endpoint distance recomputed
Synthetic cases	20	0	0	6 cases × 3 runs, plus the 230 m oracle
Sensitivity	5	0	0	only the gap setting changed
Spatial assumption	0	1	0	route transfers only at OSM junctions
External evidence	2	1	0	satellite imagery at the 3 flagged sites

What the checks caught

FAIL: a 2D crossing at a rail bridge

The route moves from way 891041250 onto bridge-tagged way 239559205 at (−3.993745, 39.54626); neither line has a vertex there. Imagery shows an **overpass**: the route cannot switch lines here.

A bug in the agent’s own checker

The connector-limit check hard-coded 50 m, so the 250 m rerun was checked against the wrong limit. It now reads the run’s setting.

A wrong claim in the Week 3 answers

The write-up said rail used a gap connector. `connectors.gpkg` shows all four used connectors are endpoint connectors (rail: 194.8 m and 372.7 m). Corrected.

A justified failure: power

`endpoint_too_far` confirmed independently: start 1,839.2 m and end 995.1 m from the nearest mapped power line, against a 500 m limit.

2026-09-27


Verifying an agentic GIS routing workflow

What the checks caught

About 45 s. First, the one failure: the audit looks for places where the route changes from one OSM line to another at a point that is not a shared vertex, which is exactly where "2D crossings connect" can be wrong. It found one, on a bridge-tagged rail line. Second, the agent’s own Week 3 discussion claimed rail used a gap connector; the saved connectors file shows all four are endpoint connectors, so the prose was wrong though the numbers were right. Third, the built-in check hard-coded 50 m and would have mis-checked the 250 m rerun. Fourth, a justified no-route result: power really is 1.8 km from the start. A successful run alone would have hidden the first three.

What the checks caught	
FAIL: a 2D crossing at a rail bridge The route moves from way 891041250 onto bridge-tagged way 239559205 at (−3.993745, 39.54626); neither line has a vertex there. Imagery shows an overpass : the route cannot switch lines here.	A wrong claim in the Week 3 answers The write-up said rail used a gap connector. <code>connectors.gpkg</code> shows all four used connectors are endpoint connectors (rail: 194.8 m and 372.7 m). Corrected.
A bug in the agent’s own checker The connector-limit check hard-coded 50 m, so the 250 m rerun was checked against the wrong limit. It now reads the run’s setting.	A justified failure: power <code>endpoint_too_far</code> confirmed independently: start 1,839.2 m and end 995.1 m from the nearest mapped power line, against a 500 m limit.

External evidence: the flagged crossing is an overpass



Claim tested in imagery	Result
Rail route switches lines at grade at the flagged crossing	FAIL overpass
No rail closer than the 195 m start connector	PASS
No rail closer than the 373 m end connector	PASS

Source: Google Maps satellite (Maxar, Airbus; ©2026), checked 27 Sep 2026; independent of OSM. Picture: Esri imagery with the model's lines.

Both OSM ways are the metre-gauge main line, and Google labels the spot “Mazeras Railway Spiral”: the line loops over itself. The audit’s FAIL was a real error.

Verifying an agentic GIS routing workflow

2026-09-27

External evidence: the flagged crossing is an overpass

About 25 s. I checked the three sites the audit flagged in satellite imagery, which is independent of OpenStreetMap. The crossing where the rail route switches lines is an overpass. Both lines are the old metre-gauge railway, and Google labels the spot the Mazeras Railway Spiral: the line loops round and crosses over itself, so there is no junction, and the rail route as computed cannot actually be followed. The two endpoint connectors hold up: there is no closer track than the mapped line, so they really are new construction. This is one image source at one date.



Sensitivity: changing only the gap tolerance changes no route outcome

Gap tolerance	Roads	Power	Rail	Power parts*	Ranking
25 m	ok, 22,897.6 m	endpoint_too_far	ok, 25,902.4 m	6	baseline < roads < rail
50 m (original)	ok, 22,897.6 m	endpoint_too_far	ok, 25,902.4 m	3	baseline < roads < rail
250 m	ok, 22,897.6 m	endpoint_too_far	ok, 25,902.4 m	1	baseline < roads < rail

*Power connected components after gap repair (9 before). Runs: sensitivity_gap_25m, default, gap_tolerance_250m.

Why nothing moves. No route uses a gap connector, and power fails on the 1,839 m endpoint distance, which no gap setting can fix. Connector counts and lengths are identical in all three runs.

What this does not show. One setting at three values is not global robustness (Saltelli et al., 2019). A common rate of USD 5, 10 or 20/m scales costs but keeps the same order.

Verifying an agentic GIS routing workflow

2026-09-27

Sensitivity: changing only the gap tolerance changes no route outcome

Sensitivity: changing only the gap tolerance changes no route outcome

Gap tolerance	Roads	Power	Rail	Power parts*	Ranking
25 m	ok, 22,897.6 m	endpoint_too_far	ok, 25,902.4 m	6	baseline < roads < rail
50 m (original)	ok, 22,897.6 m	endpoint_too_far	ok, 25,902.4 m	3	baseline < roads < rail
250 m	ok, 22,897.6 m	endpoint_too_far	ok, 25,902.4 m	1	baseline < roads < rail

*Power connected components after gap repair (9 before). Runs: sensitivity_gap_25m, default, gap_tolerance_250m.

Why nothing moves. No route uses a gap connector, and power fails on the 1,839 m endpoint distance, which no gap setting can fix. Connector counts and lengths are identical in all three runs.

What this does not show. One setting at three values is not global robustness (Saltelli et al., 2019). A common rate of USD 5, 10 or 20/m scales costs but keeps the same order.

Conclusions: what I trust and what stays unresolved

Conclusion	Label (o4_o2 §16)	Evidence
Lengths, endpoints, containment and costs are computed correctly	VERIFIED	R-*, C-*
Connectors within limits, counted once; all are endpoint connectors	VERIFIED	R-*-conn*
Power unavailable because its start is 1,839 m from mapped lines	VERIFIED; a real line there: UNRESOLVED	U-power
Ranking baseline < roads < rail holds for a 25–250 m gap tolerance	not SENSITIVE to this one setting	S-*
The rail route is continuous through the crossing	CONTRADICTED: imagery shows an overpass	V-crossings, E-crossing
Rail endpoint connectors are new construction, not unmapped track	EXTERNALLY SUPPORTED (one source, one date)	E-connector 3, 4
Access rights, permits, feasibility and real costs	UNRESOLVED	none in data

So the rail route as computed is **not valid**: the agent should forbid that transfer and re-route (or report no_path), not relax a threshold. Roads stays the route to investigate.

2026-09-27

Verifying an agentic GIS routing workflow

Conclusions: what I trust and what stays unresolved

About 35 s. What I trust: the computations are verified, and the power failure is real under this data. The ranking does not depend on the gap tolerance. What the imagery showed: the rail crossing is an overpass, so the rail route jumps between two lines that do not connect; it is verified as a computation but invalid as a route. Both rail endpoint connectors hold up: no closer track than the mapped line. That is one image source at one date. Nothing here establishes access rights, permits or real costs; proximity is not permission, and 10 dollars per metre is a classroom value. Verified, not yet validated.

Conclusions: what I trust and what stays unresolved		
Conclusion	Label (o4_o2 §16)	Evidence
Lengths, endpoints, containment and costs are computed correctly	VERIFIED	R-*, C-*
Connectors within limits, counted once; all are endpoint connectors	VERIFIED	R-*-conn*
Power unavailable because its start is 1,839 m from mapped lines	VERIFIED; a real line there: UNRESOLVED	U-power
Ranking baseline < roads < rail holds for a 25–250 m gap tolerance	not SENSITIVE to this one setting	S-*
The rail route is continuous through the crossing	CONTRADICTED: imagery shows an overpass	V-crossings, E-crossing
Rail endpoint connectors are new construction, not unmapped track	EXTERNALLY SUPPORTED (one source, one date)	E-connector 3, 4
Access rights, permits, feasibility and real costs	UNRESOLVED	none in data
So the rail route as computed is not valid : the agent should forbid that transfer and re-route (or report no_path), not relax a threshold. Roads stays the route to investigate.		

Data, tools and readings

- ▶ Roads, rail, power, settlements: © OpenStreetMap contributors, ODbL (saved in o3_o1).
- ▶ Basemaps: Esri World Light Gray Canvas and World Street Map (Esri, HERE, Garmin, © OpenStreetMap contributors).
- ▶ GGS 662 notebooks o3_o1, o3_o2, o4_o1, o4_o2 and AGENTS.md. Python: geopandas, shapely, networkx, pandas, pyproj.
- ▶ Oreskes, Shrader-Frechette & Belitz (1994), *Science*, doi:10.1126/science.263.5147.641
- ▶ Haklay (2010), *Environment and Planning B*, doi:10.1068/b35o97
- ▶ Saltelli et al. (2019), *Environmental Modelling & Software*, doi:10.1016/j.envsoft.2019.01.012

AI-use statement

Tool: Claude Code (Anthropic, Claude Opus 5.5).

How it contributed: general coding assistance, writing and debugging the Week 3 routing script, the Week 4 audit, and the plotting and export scripts, and running the reruns; drafting the expected-checks table; and generating these slides.

My role: I directed the analysis, checked the three flagged sites in satellite imagery myself, reviewed the results and directed edits to the slides. I remain responsible for understanding, testing and validating all submitted material. Because the same tool wrote the routing code and the audit, I treat the audit as verification, not independent validation.

2026-09-27

Sources and AI use

About 10 s. Data from OpenStreetMap contributors under ODbL, with Esri basemaps. Readings are the three from o4_o2 that shaped the audit. AI use: Claude Code wrote and ran the code and generated the slides; I directed it, did the imagery check and reviewed the results.

Sources and AI use	
Data, tools and readings ▶ Roads, rail, power, settlements: © OpenStreetMap contributors, ODbL (saved in o3_o1). ▶ Basemaps: Esri World Light Gray Canvas and World Street Map (Esri, HERE, Garmin, © OpenStreetMap contributors). ▶ GGS 662 notebooks o3_o1, o3_o2, o4_o1, o4_o2 and AGENTS.md. Python: geopandas, shapely, networkx, pandas, pyproj. ▶ Oreskes, Shrader-Frechette & Belitz (1994), <i>Science</i>, doi:10.1126/science.263.5147.641 ▶ Haklay (2010), <i>Environment and Planning B</i>, doi:10.1068/b35o97 ▶ Saltelli et al. (2019), <i>Environmental Modelling & Software</i>, doi:10.1016/j.envsoft.2019.01.012	AI-use statement Tool: Claude Code (Anthropic, Claude Opus 5.5). How it contributed: general coding assistance, writing and debugging the Week 3 routing script, the Week 4 audit, and the plotting and export scripts, and running the reruns; drafting the expected-checks table, and generating these slides. My role: I directed the analysis, checked the three flagged sites in satellite imagery myself, reviewed the results and directed edits to the slides. I remain responsible for understanding, testing and validating all submitted material. Because the same tool wrote the routing code and the audit, I treat the audit as verification, not independent validation.

Thank you Questions?

Supporting slides follow: full check table, synthetic cases and evidence files


2026-09-27

Verifying an agentic GIS routing workflow

Thank you
Questions?

Supporting slides follow: full check table, synthetic cases and evidence files

Supporting: each candidate route against the baseline




Roads follows the baseline almost exactly (+0.8 m). Power has no route: its start is 1,839 m from the nearest mapped line. Rail is 3,006 m longer, and imagery shows it is **not followable** at the Mazeras crossing. Source: route_side_by_side.png, 50 m run.

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: each candidate route against the baseline



Supporting: expected checks, written as specifications (04_02 §13)

Check	Expected / tolerance	Failure would mean	Audit IDs
Endpoints	Easterly start, westerly end within 0.01 m	Wrong problem solved	R-*-ends
Continuity, validity	One valid LineString, no gaps or loops	Length and cost meaningless	R-*-valid, -simple
Containment	Inside baseline bounds + 3,000 m	Data outside allowed environment	R-*-contain
Metres	EPSG:32737; edge sum within 0.01 m; UTM vs geodesic <0.5 %; ≥ straight line	CRS or double-count error	R-*-len, -units
Connectors	Gap ≤ run tolerance; endpoint ≤ 500 m; counted once; baseline 0	Invented or double-counted construction	R-*-conn*, C-base
Cost	length × USD 10 within USD 0.01, one rate	Wrong ranking	R-*-cost, C-rate
Unavailable	Named status, blank length/cost, no geometry	Failure hidden as a route	C-unavail, U-power
Synthetic	5 AGENTS.md cases; 80+120+30 m → 230 m, USD 2,300	Logic bug hidden by real data	T-*
Sensitivity	Only gap tolerance changes (25 m); changes reported	Uncontrolled or forced result	S-*
Junctions	Transfers only at shared OSM vertices	Bridge or tunnel treated as a junction	V-crossings

Geometry cannot establish: buildability, access rights or permits; a 2D crossing may be a bridge (needs imagery, E-*). Full table: Claude/04_Project/expected_checks.ipynb.

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: expected checks, written as specifications (04_02 §13)

Supporting: expected checks, written as specifications (04_02 §13)			
Check	Expected / tolerance	Failure would mean	Audit IDs
Endpoints	Easterly start, westerly end within 0.01 m	Wrong problem solved	R-*-ends
Continuity, validity	One valid LineString, no gaps or loops	Length and cost meaningless	R-*-valid, -simple
Containment	Inside baseline bounds + 3,000 m	Data outside allowed environment	R-*-contain
Metres	EPSG:32737; edge sum within 0.01 m; UTM vs geodesic <0.5 %; ≥ straight line	CRS or double-count error	R-*-len, -units
Connectors	Gap ≤ run tolerance; endpoint ≤ 500 m; counted once; baseline 0	Invented or double-counted construction	R-*-conn*, C-base
Cost	length × USD 10 within USD 0.01, one rate	Wrong ranking	R-*-cost, C-rate
Unavailable	Named status, blank length/cost, no geometry	Failure hidden as a route	C-unavail, U-power
Synthetic	5 AGENTS.md cases; 80+120+30 m → 230 m, USD 2,300	Logic bug hidden by real data	T-*
Sensitivity	Only gap tolerance changes (25 m); changes reported	Uncontrolled or forced result	S-*
Junctions	Transfers only at shared OSM vertices	Bridge or tunnel treated as a junction	V-crossings
Geometry cannot establish: buildability, access rights or permits; a 2D crossing may be a bridge (needs imagery, E-*). Full table: Claude/04_Project/expected_checks.ipynb.			

Supporting: per-route checks (all PASS)

Check (tolerance)	Baseline	Roads	Rail
Valid, nonempty single LineString	PASS	PASS	PASS
Endpoints match baseline (≤ 0.01 m)	0.0000 / 0.0000 m	0.0000 / 0.0000 m	0.0000 / 0.0000 m
No loops or self-crossing (is_simple)	True	True	True
Inside study rectangle	True	True	True
Geometry vs traversed-edge sum (≤ 0.01 m)	0.000000 m	0.009833 m	0.007000 m
Metres not degrees (UTM vs geodesic, $< 0.1\%$)	−0.036%	−0.036%	−0.036%
Not shorter than straight line	1.111×	1.111×	1.257×
Cost = length × USD 10/m	228,967.56	228,975.93	259,023.67
Connectors counted once (count, total)	0, 0 m	2, 1.3139 m	2, 567.5396 m
Each connector ≤ 500 m, inside, on route, at endpoint	–	0.59, 0.72 m	194.85, 372.69 m
Non-connector length on mapped lines (± 1 m)	–	22,896.4 m	25,334.9 m

Source: validation_table.csv (R-* checks). Note: roads edge-sum difference (0.0098 m) passes but sits close to the 0.01 m tolerance, consistent with 1 cm node rounding.

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: per-route checks (all PASS)

Check (tolerance)	Baseline	Roads	Rail
Valid, nonempty single LineString	PASS	PASS	PASS
Endpoints match baseline (≤ 0.01 m)	0.0000 / 0.0000 m	0.0000 / 0.0000 m	0.0000 / 0.0000 m
No loops or self-crossing (is_simple)	True	True	True
Inside study rectangle	True	True	True
Geometry vs traversed-edge sum (≤ 0.01 m)	0.000000 m	0.009833 m	0.007000 m
Metres not degrees (UTM vs geodesic, $< 0.1\%$)	−0.036%	−0.036%	−0.036%
Not shorter than straight line	1.111×	1.111×	1.257×
Cost = length × USD 10/m	228,967.56	228,975.93	259,023.67
Connectors counted once (count, total)	0, 0 m	2, 1.3139 m	2, 567.5396 m
Each connector ≤ 500 m, inside, on route, at endpoint	–	0.59, 0.72 m	194.85, 372.69 m
Non-connector length on mapped lines (± 1 m)	–	22,896.4 m	25,334.9 m

Source: validation_table.csv (R-* checks). Note: roads edge-sum difference (0.0098 m) passes but sits close to the 0.01 m tolerance, consistent with 1 cm node rounding.

Supporting: provenance, data QA and comparison table

Check ID	Observed	Status
P-settings	Settings match spec: gap 50 m, endpoint 500 m, margin 3,000 m, EPSG:32737	PASS
D-roads	4,742 features, EPSG:4326, 0 missing / empty / invalid, all intersect study area	PASS
D-power	15 features, 0 missing / empty / invalid	PASS
D-rail	77 features, 0 missing / empty / invalid	PASS
D-baseline	1 LineString, valid	PASS
D-settlements	30 features (29 points, 1 polygon), map context only	PASS
D-roads-dup	2,292 of 4,742 edges are reversed OSMnx duplicates; removed before routing	PASS
C-rows	baseline, roads, power, rail: one row each	PASS
C-cols	nine required CSV columns	PASS
C-base	baseline connector values 0, 0, 0	PASS
C-unavail	power: endpoint_too_far, blank length and cost (not zero)	PASS
C-gpkg	routes.gpkg holds baseline, roads, rail	PASS

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: provenance, data QA and comparison table

Check ID	Observed	Status
P-settings	Settings match spec: gap 50m, endpoint 500m, margin 3,000m, EPSG:32737	PASS
D-roads	4,742 features, EPSG:4326, 0 missing / empty / invalid, all intersect study area	PASS
D-power	15 features, 0 missing / empty / invalid	PASS
D-rail	77 features, 0 missing / empty / invalid	PASS
D-baseline	1 LineString, valid	PASS
D-settlements	30 features (29 points, 1 polygon), map context only	PASS
D-roads-dup	2,292 of 4,742 edges are reversed OSMnx duplicates; removed before routing	PASS
C-rows	baseline, roads, power, rail: one row each	PASS
C-cols	nine required CSV columns	PASS
C-base	baseline connector values 0, 0, 0	PASS
C-unavail	power: endpoint_too_far, blank length and cost (not zero)	PASS
C-gpkg	routes.gpkg holds baseline, roads, rail	PASS

Supporting: unavailable route, spatial assumption, sensitivity, external evidence

Check ID	Observed	Status
C-rate	one common rate, USD 10/m	PASS
U-power	start 1,839.2 m, end 995.1 m from nearest power line (>500 m)	PASS
V-crossings	1 transfer at a non-shared vertex; involves a bridge-tagged way	FAIL
S-25m-ctrl	only “Gap connector limit” changed, 50 → 25 m	PASS
S-25m-out	no differences; ranking unchanged; power parts 6	PASS
S-250m-ctrl	only “Gap connector limit” changed, 50 → 250 m	PASS
S-250m-out	no differences; ranking unchanged; power parts 1	PASS
S-rate	same order at USD 5, 10, 20/m	PASS
E-rail connector 3	195 m connector (−4.008849, 39.591099): no track closer than the mapped line	PASS
E-rail connector 4	373 m connector (−3.864413, 39.472359): mapped line is the closest rail	PASS
E-rail crossing 1	crossing (−3.993745, 39.54626): overpass, not a junction	FAIL

Source: Google Maps satellite imagery (Maxar, Airbus; ©2026, credit year, not exact capture date), independent of OSM. Recorded in external_evidence.csv; site maps site_*.png.

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: unavailable route, spatial assumption, sensitivity, external evidence

Supporting: unavailable route, spatial assumption, sensitivity, external evidence		
Check ID	Observed	Status
C-rate	one common rate, USD 10/m	PASS
U-power	start 1,839.2 m, end 995.1 m from nearest power line (>500 m)	PASS
V-crossings	1 transfer at a non-shared vertex; involves a bridge-tagged way	FAIL
S-25m-ctrl	only “Gap connector limit” changed, 50 → 25 m	PASS
S-25m-out	no differences; ranking unchanged; power parts 6	PASS
S-250m-ctrl	only “Gap connector limit” changed, 50 → 250 m	PASS
S-250m-out	no differences; ranking unchanged; power parts 1	PASS
S-rate	same order at USD 5, 10, 20/m	PASS
E-rail connector 3	195 m connector (−4.008849, 39.591099): no track closer than the mapped line	PASS
E-rail connector 4	373 m connector (−3.864413, 39.472359): mapped line is the closest rail	PASS
E-rail crossing 1	crossing (−3.993745, 39.54626): overpass, not a junction	FAIL
Source: Google Maps satellite imagery (Maxar, Airbus; ©2026, credit year, not exact capture date), independent of OSM. Recorded in external_evidence.csv; site maps site_*.png.		

Supporting: synthetic cases and known-answer oracle

Case (AGENTS.md)	Expected	50 m run	25 m run	250 m run
Already connected: components	1	1 PASS	1 PASS	1 PASS
Gap below threshold: components after repair	1	1 PASS	1 PASS	1 PASS
Gap above threshold: components after repair	2	2 PASS	2 PASS	2 PASS
Endpoint on a line interior: status	ok	ok PASS	ok PASS	ok PASS
Disconnected endpoints: status	no_path	no_path PASS	no_path PASS	no_path PASS
Connector counted once: route length (10 m + gap + 20 m)	= expected	55.0 m PASS	42.5 m PASS	155.0 m PASS
Oracle: 80 m + 120 m + 30 m gap at USD 10/m	230 m, USD 2,300	230 m, USD 2,300 PASS		
Oracle fault: gap double-counted	flagged	260 m $\neq$ 230 m flagged PASS		

The gap in the connector case is half the run’s tolerance, so the expected length differs per run. Source: routing_report.md in each run folder; audit IDs T-*.

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: synthetic cases and known-answer oracle

Case (AGENTS.md)	Expected	50 m run	25 m run	250 m run
Already connected: components	1	1 PASS	1 PASS	1 PASS
Gap below threshold: components after repair	1	1 PASS	1 PASS	1 PASS
Gap above threshold: components after repair	2	2 PASS	2 PASS	2 PASS
Endpoint on a line interior: status	ok	ok PASS	ok PASS	ok PASS
Disconnected endpoints: status	no_path	no_path PASS	no_path PASS	no_path PASS
Connector counted once: route length (10 m + gap + 20 m)	= expected	55.0 m PASS	42.5 m PASS	155.0 m PASS
Oracle: 80 m + 120 m + 30 m gap at USD 10/m	230 m, USD 2,300	230 m, USD 2,300 PASS		
Oracle fault: gap double-counted	flagged	260 m $\neq$ 230 m flagged PASS		

The gap in the connector case is half the run's tolerance, so the expected length differs per run. Source: routing_report.md in each run folder; audit IDs T-*

Supporting: evidence files, run commands and versions

Run commands (every argument has a default)

- ▶ `python Claude/03_Project/week3_routing.py → outputs/default/`
- ▶ `python week3_routing.py --gap-tolerance-m 25 --out outputs/sensitivity_gap_25m`
- ▶ `python week3_routing.py --gap-tolerance-m 250 --out outputs/gap_tolerance_250m`
- ▶ `python Claude/04_Project/week4_audit.py → outputs/week4_validation/`

Evidence files

- ▶ `outputs/week3_routing/` and each rerun subfolder: `route_comparison.csv`, `routes.gpkg`, `connectors.gpkg`, `route_map.png`, `routing_report.md`
- ▶ `outputs/week4_validation/`: `week4_audit.py`, `validation_table.csv`, `validation_report.md`, `validation_map.png`, `inferred_crossings.gpkg`, `external_evidence.csv`

Versions: Python 3.12.2; geopandas 1.1.4; shapely 2.1.2; pyproj 3.8.0; networkx 3.6.1; pandas 2.3.3.

Output location: as O4_O2 specifies, under the repo root: `outputs/week3_routing/` (original run, plus `sensitivity_gap_25m/` and `gap_tolerance_250m/`) and `outputs/week4_validation/`. Copied from the working folders under Claude/ by Claude/04_Project/export_evidence.py.

Verifying an agentic GIS routing workflow

2026-09-27

Supporting: evidence files, run commands and versions

Supporting: evidence files, run commands and versions

Run commands (every argument has a default)

- ▶ `python Claude/03_Project/week3_routing.py → outputs/default/`
- ▶ `python week3_routing.py --gap-tolerance-m 25 --out outputs/sensitivity_gap_25m`
- ▶ `python week3_routing.py --gap-tolerance-m 250 --out outputs/gap_tolerance_250m`
- ▶ `python Claude/04_Project/week4_audit.py → outputs/week4_validation/`

Evidence files

- ▶ `outputs/week3_routing/` and each rerun subfolder: `route_comparison.csv`, `routes.gpkg`, `connectors.gpkg`, `route_map.png`, `routing_report.md`
- ▶ `outputs/week4_validation/`: `week4_audit.py`, `validation_table.csv`, `validation_report.md`, `validation_map.png`, `inferred_crossings.gpkg`, `external_evidence.csv`

Versions: Python 3.12.2; geopandas 1.1.4; shapely 2.1.2; pyproj 3.8.0; networkx 3.6.1; pandas 2.3.3.

Output location: as O4_O2 specifies, under the repo root: `outputs/week3_routing/` (original run, plus `sensitivity_gap_25m/` and `gap_tolerance_250m/`) and `outputs/week4_validation/`. Copied from the working folders under Claude/ by Claude/04_Project/export_evidence.py.